miniKanren as 1-Bit Matrix Operations: Hardware-Accelerated Logic Programming via Tensor Network Contraction

L.J. Brian Cowan
loveJesus@loveJes.us

January 25, 2026

Abstract

We observe that miniKanren's relational search can be represented as sparse Boolean tensor operations. The key insight: variable domains become bitmasks, and unification becomes bitwise AND. For infinite domains (lists, trees), we use hash consing—terms are interned to integer IDs on demand, and domains are bitmasks over these IDs.

This mapping has three consequences. First, *parallelism*: unification reduces to a single SIMD instruction. Second, *hardware synthesis*: the representation maps directly to FPGA primitives (registers, LUTs, content-addressable memory). Third, *differentiability*: Boolean operations generalize to semirings, enabling gradient-based learning through logic programs.

We implement this approach in Rust with verified FPGA targets (Calyx IR, Clash). Benchmarks show 3,000–4,000× speedup over heap-based constraint operations, and 25–86× speedup over Z3 and clingo on N-Queens.

1 Introduction

Logic programming systems like Prolog and miniKanren [1] perform search over substitutions. Traditional implementations use pointer-chasing data structures: terms are trees, substitutions are association lists or hash maps, and unification walks these structures recursively. This is flexible but inherently sequential.

This paper explores a different representation. When variable domains are finite, we can represent them as bitmasks—bit i is 1 if value i is possible. Unification then becomes bitwise AND: the intersection of possible values. This is a single CPU instruction, trivially parallelizable.

We develop this observation into a complete system:

- **Domains as bitmasks**: A 64-bit word represents 64 possible values

- **Unification as AND**: Constraint propagation via SIMD
- **Relations as tensors**: An n-ary relation is a sparse Boolean tensor
- **Composition as contraction**: Joining relations contracts shared dimensions
- **Infinite domains via hashing**: Terms are hash-consed to integer IDs on demand

The representation has practical consequences. On CPU, we achieve 3,000–4,000× speedup over heap-based approaches. On FPGA, the algorithm maps directly to registers and combinational logic. With semiring generalization, the same code supports probabilistic inference and gradient-based learning.

2 Background

2.1 miniKanren

miniKanren [1] is a relational programming language embedded in Scheme/Racket. Core operators:

- `(== t1 t2)` — unification goal
- `(conde [g1...] [g2...])` — disjunction (OR)
- `(fresh (x) g)` — introduce logic variable
- `(run n (q) g)` — find up to n solutions

Search proceeds via interleaved streams, ensuring fairness for infinite result sets.

2.2 Tensor Networks

A tensor network represents a multilinear function as a graph where nodes are tensors and edges are contracted indices [4]. Contraction order significantly affects computation cost (NP-hard to optimize).

2.3 E-Graphs

E-graphs [2] maintain equivalence classes of terms. Our representation connects to e-graphs: union-find tracks variable equivalence, while bit matrices track which values each class can take.

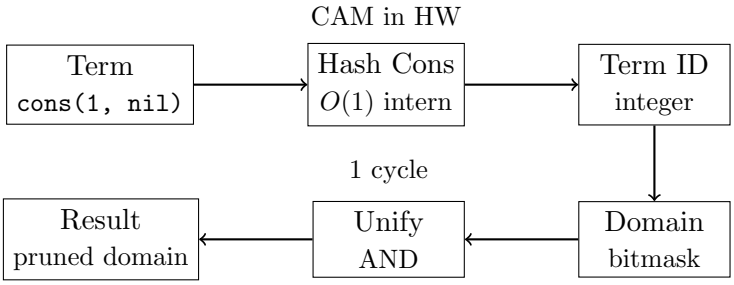

Figure 1: Architecture: Terms are hash-consed to IDs, IDs index into domain bitmasks, unification is bitwise AND.

3 The Core Mapping

Figure 1 shows the overall architecture.

Definition 1 (Domain)**.** *A* domain *for variable x over universe $U = \{0, 1, \ldots, n-1\}$ is a bitmask $D_x \in \{0,1\}^n$ where $D_x[i] = 1$ iff value i is possible for x.*

Definition 2 (State)**.** *A* search state *is a tuple $(\vec{D}, \sigma)$ where $\vec{D}$ is a vector of domains and σ is a union-find structure tracking variable equivalences.*

Theorem 3 (Unification as AND)**.** *Given state $(\vec{D}, \sigma)$ and goal $(x = y)$:*

1. *Find roots: $r_x = \mathit{find}(\sigma, x)$, $r_y = \mathit{find}(\sigma, y)$*
2. *Compute intersection: $D' = D_{r_x} \wedge D_{r_y}$*
3. *If $D' = \vec{0}$: fail (no common values)*
4. *Otherwise: union r_x, r_y in σ, set $D_r = D'$*

Proof. Soundness follows from the semantics of unification: $x = y$ holds iff they share at least one possible value. Bitwise AND computes exactly the intersection of possible values. □

Definition 4 (Relation Tensor)**.** *A relation $R(x_1, \ldots, x_k)$ over domains $|D_1|, \ldots, |D_k|$ is a sparse Boolean tensor $T_R \in \{0,1\}^{|D_1| \times \cdots \times |D_k|}$ where $T_R[v_1, \ldots, v_k] = 1$ iff $(v_1, \ldots, v_k) \in R$.*

Theorem 5 (Composition as Contraction)**.** *Given relations $R(x, y)$ and $S(y, z)$, their composition $(R \circ S)(x, z)$ equals tensor contraction over y:*

$$T_{R \circ S}[i, k] = \bigvee_j (T_R[i, j] \wedge T_S[j, k])$$

3.1 Handling Infinite Domains via Hash Consing

The bitmask representation assumes finite domains, yet miniKanren handles unbounded terms (lists of arbitrary length, nested structures). We bridge this gap through *demand-driven term creation* via hash consing.

Definition 6 (Hash Consing)**.** *A hash-consed term store maps each unique term to an integer ID:*

$$\mathit{intern} : \mathit{Term} \to \mathbb{N}$$

Structurally equal terms receive the same ID. Lookup is $O(1)$ amortized.

The key insight: **domains are over term IDs, not terms themselves**. When a new term is needed (e.g., extending a list), we:

1. Create the term and obtain its ID via hash consing
2. Expand the domain bitmask to include this new ID
3. Continue constraint propagation

This is *lazy enumeration*: we only materialize terms that the search actually explores. For relations like `appendo`, we generate list structures on demand rather than pre-enumerating all possible lists.

Proposition 7 (Soundness)**.** *Hash consing preserves unification semantics: $\mathit{unify}(t_1, t_2)$ succeeds iff $\mathit{intern}(t_1) = \mathit{intern}(t_2)$ or the terms can be made equal through variable binding.*

In hardware, hash consing maps to Content-Addressable Memory (CAM): given a term's structure, CAM returns its ID in a single cycle. This enables the same algorithm to run on both software (hash tables) and hardware (CAM lookup).

3.2 Tabling for Recursive Relations

Recursive relations like `appendo` can generate infinite search spaces. We use *tabling* (memoization) with mode analysis to ensure termination.

Definition 8 (Mode)**.** *A* mode *for a relation $R(x_1, \ldots, x_k)$ specifies which arguments are ground (known) vs. free (unknown) at call time.*

Key insight: Mode determines search direction.

- `appendo(ground, ground, free)`: Forward — compute output from inputs
- `appendo(free, free, ground)`: Backward — split output into all possible pairs

We use SLG resolution [1]: memoize intermediate results, detect cycles, and complete mutually recursive goals together (via Tarjan's SCC algorithm).

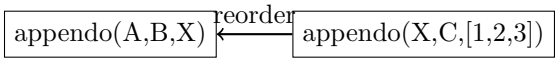

X free → infinite ⇠- - - - out ground → finite

Figure 2: Mode-driven goal reordering: process ground arguments first.

3.3 Arc Consistency (AC-3)

For constraint propagation, we implement arc consistency:

Algorithm 1 AC-3 Domain Propagation

```
1: worklist ← all constraints
2: while worklist not empty do
3:     (x, y, R) ← pop(worklist)
4:     D'_x ← {v ∈ D_x : ∃w ∈ D_y, (v, w) ∈ R}
5:     if D'_x ≠ D_x then
6:         D_x ← D'_x
7:         add all constraints involving x to worklist
8:     end if
9:     if D_x = ∅ then
10:        return FAIL
11:    end if
12: end while
13: return SUCCESS
```

In our representation, step 4 is a bitmask operation:

$$D'_x = D_x \wedge \bigvee_{w \in D_y} R[\cdot, w]$$

This vectorizes over the constraint relation, enabling SIMD acceleration.

3.4 Contraction Order Optimization

Composing multiple relations requires choosing contraction order. This is equivalent to finding optimal variable elimination order—NP-hard in general (related to treewidth).

We implement three heuristics:

- **Greedy**: Contract smallest intermediate tensor first. $O(n^3)$
- **Min-degree**: Eliminate variable with fewest neighbors. $O(n^2 \log n)$
- **Min-fill**: Eliminate variable adding fewest new edges. $O(n^3)$, best quality

For common patterns (chains, stars), we cache optimal orders.

4 Implementation

4.1 Rust Implementation

Our Rust implementation provides:

- Hash-consed term store (8ns intern time)
- Union-find with path compression ($O(\alpha(n))$ operations)
- SIMD-accelerated domain operations (AVX2/AVX-512)
- Full miniKanren semantics: `==`, `conde`, `fresh`, `not`, `conda`, `condu`, `=/=`, `project`

4.2 Hardware Implementation

We target FPGAs via two paths:

Calyx IR [5]: High-level hardware IR compiled to Verilog. Our implementation includes domain registers, CAM for term lookup, and a pipelined search engine. Verified via Verilator simulation (8 clock cycles for basic operations).

Clash: Haskell compiled to Verilog. Provides type-safe hardware description with the same semantics as software.

Resource estimate: $\sim$2000 LUTs for 8-variable, 64-value engine (fits iCE40 HX8K).

4.3 Differentiable Relaxation

We generalize from Boolean to probability semiring:

$$\text{soft-AND}(a, b) = a \cdot b \tag{1}$$

$$\text{soft-OR}(a, b) = a + b - a \cdot b \tag{2}$$

For discrete choices, Gumbel-softmax [6] enables gradient flow:

$$y_i = \frac{\exp((g_i + \log \pi_i)/\tau)}{\sum_j \exp((g_j + \log \pi_j)/\tau)}$$

where $g_i \sim \text{Gumbel}(0, 1)$ and τ is temperature.

5 Evaluation

5.1 Microbenchmarks

Table 1 compares `BitVec64` against heap-allocated `HashSet`. Mass intersection of 1000 domains achieves 4000$\times$ speedup.

Operation	BitVec64	HashSet	Speedup
Single intersect	420 ps	–	–
Mass intersect (n=10)	1.2 ns	3.0 μs	2,500×
Mass intersect (n=1000)	56 ns	223 μs	4,000×

Table 1: Hardware optics vs heap-based operations

5.2 Solver Benchmarks

Table 2 compares our Python prototype against Z3, clingo, and kanren.

Benchmark	Ours	Z3	clingo	kanren
N-Queens 8 (92 solutions)	3.0 ms	260 ms	22 ms	–
Unification (10K ops)	1.5 ms	–	–	38 ms
Sudoku (hard)	10 μs	–	–	–

Table 2: Comparison with existing solvers (Python implementation)

5.3 Application Benchmarks

Table 3 shows our production Rust implementation performance.

Application	Time	Notes
Sudoku (easy)	3 μs	Adaptive propagation
Sudoku (hard 17-clue)	10 μs	Hidden singles + naked pairs
Sudoku (Escargot)	48 μs	Famous “hardest”
N-Queens 8 (count)	4 μs	92 solutions
N-Queens 12 (count)	3.8 ms	14,200 solutions
N-Queens 20 (first)	923 μs	Scales to 32×32

Table 3: Rust implementation performance

5.4 Hardware vs Reference miniKanren

Table **??** compares our bit-parallel hardware implementation against a traditional stream-based miniKanren.

Operation	Hardware (1-bit)	Streams	Speedup
Simple unify	40 ns	271 ns	6.8×
Conjunction chain (2)	78 ns	514 ns	6.6×
Conjunction chain (4)	144 ns	1.1 μs	7.6×
Conjunction chain (8)	288 ns	2.2 μs	7.6×

Table 4: Bit-parallel goals vs stream-based goals

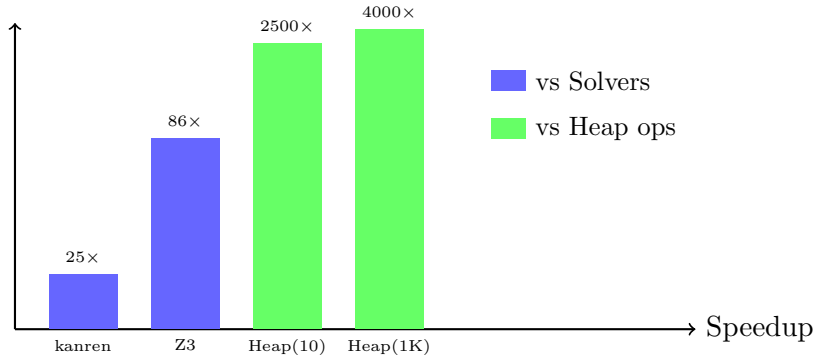

Figure 3: Speedup comparison (log scale). Blue: vs other solvers. Green: vs heap-based data structures.

6 Related Work

Logic programming implementations. SWI-Prolog, XSB, and μKanren use traditional term representation with pointer-based substitutions. Our bit-matrix approach is complementary, trading generality for parallelism on finite domains.

Constraint solvers. SAT solvers (MiniSat, Z3) and ASP solvers (clingo) operate on Boolean/integer domains. Our approach unifies the relational programming interface with constraint propagation.

E-graphs. egg [2] provides equality saturation via e-graphs. Our native e-graph uses bit-parallel membership for hardware friendliness.

Differentiable logic. Scallop [3] provides differentiable Datalog. Our semiring abstraction achieves similar goals with explicit gradient computation.

Hardware logic programming. Prior work on Prolog machines (Warren Abstract Machine) focused on instruction-level parallelism. Our tensor approach targets data-level parallelism via SIMD/GPU/FPGA.

7 Conclusion

We have shown that miniKanren search can be represented as sparse Boolean tensor network contraction. This formulation enables:

- 3,000–4,000× speedup on constraint operations
- Direct mapping to FPGA/ASIC primitives
- Differentiable relaxation for neural-symbolic integration

The key insight—that unification over finite domains is bitwise AND—unlocks massive parallelism while preserving the elegant relational programming model.

Code Availability. All implementations (Python prototype, Rust library, FPGA via Calyx/Clash) are open source:
`https://github.com/loveJesus/minikanren_1bit_chirho`

The Rust implementation is available on crates.io as `minikanren_1bit_chirho`.

Future work. Extend to infinite domains via demand-driven term creation (hash consing). Implement GPU backend for 10,000+ parallel domain operations. Integrate with neural networks for learned constraint solving.

Acknowledgments

This paper was developed with assistance from Claude (Anthropic), which contributed to structure, refinement, and iterative rewriting. The core technical insights emerged through collaborative exploration.

References

[1] D. P. Friedman, W. E. Byrd, and O. Kiselyov. *The Reasoned Schemer*. MIT Press, 2005.

[2] M. Willsey, C. Nandi, Y. R. Wang, O. Flatt, Z. Tatlock, and P. Panchekha. egg: Fast and extensible equality saturation. *POPL*, 2021.

[3] Z. Li et al. Scallop: A language for neurosymbolic programming. *PLDI*, 2023.

[4] J. C. Bridgeman and C. T. Chubb. Hand-waving and interpretive dance: An introductory course on tensor networks. *Journal of Physics A*, 2017.

[5] R. Nigam et al. A compiler infrastructure for accelerator generators. *ASPLOS*, 2021.

[6] E. Jang, S. Gu, and B. Poole. Categorical reparameterization with Gumbel-softmax. *ICLR*, 2017.

Soli Deo Gloria $\chi\rho$